

# Finite representation of real numbers

## Floating-point numbers

Dr. Ing. Rodrigo Gonzalez

`rodrazalez@frm.utn.edu.ar`

`rodrigo.gonzalez@ingenieria.uncuyo.edu.ar`

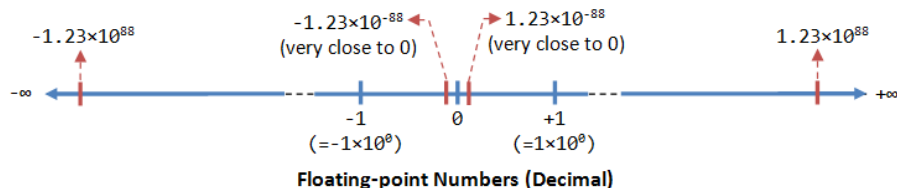


**UNCUYO**  
UNIVERSIDAD  
NACIONAL DE CUYO

# Summary

- 1 Representation
- 2 Old formats
- 3 IEEE 754 standard
  - 32-bit Single-Precision
  - Normalized Form
  - Example 1
  - Example 2
  - Example 3
  - Why auto range?
  - De-normalized Form
  - Example 4
- 4 Special values
- 5 Rounding schemes
- 6 Dynamic range
- 7 Precision
  - Fixed-point precision
  - Floating-point precision
  - Precision is not constant
  - Precision problems
  - Sum of two floating-point numbers
  - Sum of similar floating-point numbers
  - Sum of not similar floating-point numbers

A floating-point number can represent a very large or a very small value, positive and negative.



A floating-point number is typically expressed in the scientific notation in the form of

$$(-1)^S \times F \times r^E,$$

where,

- $S$ , sign bit.
- $F$ , fraction.
- $E$ , **biased** exponent.
- $r$ , certain radix.  $r = 2$  for binary;  $r = 10$  for decimal.

## IEEE Standard P754 Format

Bit	31	30	29	28	27	26	25	24	23	22	21	20	...	2	1	0
	S	$2^7$	$2^6$	$2^5$	$2^4$	$2^3$	$2^2$	$2^1$	$2^0$	$2^{-1}$	$2^{-2}$	$2^{-3}$	...	$2^{-21}$	$2^{-22}$	$2^{-23}$
Sign (s)	← Exponent (e) →									← Fraction (f) →						

## IBM Format

Bit	31	30	29	28	27	26	25	24	23	22	21	20	...	2	1	0
	S	$2^6$	$2^5$	$2^4$	$2^3$	$2^2$	$2^1$	$2^0$	$2^{-1}$	$2^{-2}$	$2^{-3}$	$2^{-4}$	...	$2^{-22}$	$2^{-23}$	$2^{-24}$
Sign (s)	← Exponent (e) →									← Fraction (f) →						

## DEC (Digital Equipment Corp.) Format

Bit	31	30	29	28	27	26	25	24	23	22	21	20	...	2	1	0
	S	$2^7$	$2^6$	$2^5$	$2^4$	$2^3$	$2^2$	$2^1$	$2^0$	$2^{-2}$	$2^{-3}$	$2^{-4}$	...	$2^{-22}$	$2^{-23}$	$2^{-24}$
Sign (s)	← Exponent (e) →									← Fraction (f) →						

## MIL-STD 1750A Format

Bit	31	30	29	...	11	10	9	8	7	6	5	4	3	2	1	0
	$2^0$	$2^{-1}$	$2^{-2}$	...	$2^{-20}$	$2^{-21}$	$2^{-22}$	$2^{-23}$	$2^7$	$2^6$	$2^5$	$2^4$	$2^3$	$2^2$	$2^1$	$2^0$
← Fraction (f) →									← Exponent (e) →							

Modern computers adopt the IEEE 754 standard for representing floating-point numbers at the FPU.

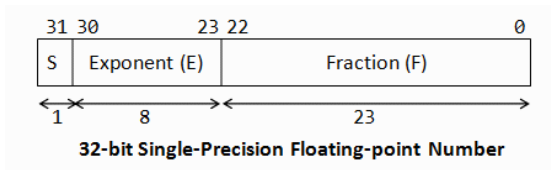
First version was published in 1985. Last version in July 2019 (IEEE 754-2019).

IEEE 754 standard defines several arithmetic formats.

Parameter	Binary formats ( $B = 2$ )				Decimal formats ( $B = 10$ )		
	Binary 16	Binary 32	Binary 64	Binary 128	Decimal 132	Decimal 164	Decimal 128
$p$ , digits	$10 + 1$	$23 + 1$	$52 + 1$	$112 + 1$	7	16	34
$e_{max}$	+15	+127	+1023	+16383	+96	+384	+16,383
$e_{min}$	-14	-126	-1022	-16382	-95	-383	-16,382
Common name	Half precision	Single precision	Double precision	Quadruple precision			

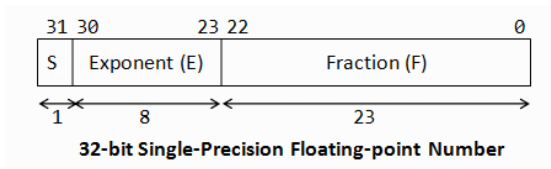
IEEE 754 standard also defines:

- Rounding rules.
- Arithmetic operations, trigonometric functions.
- Exception handling.



$$(-1)^S \times F \times r^{(E-bias)}$$

- S, sign bit. **0** for positive numbers and **1** for negative numbers.
- F, 23-bits fraction:  $[2^{-1} \ 2^{-2} \ \dots \ 2^{-23}]$
- We need to represent both positive and negative exponents.
- E, 8-bits exponent, **no sign bit**.
  - $E = [1, 254]$ ,  $bias = 127$ ;  $-126 \leq E - bias \leq 127$ .
  - $E = 0$  and  $E = 255$  are reserved.



$$(-1)^S \times F \times r^{(E-bias)}$$

- Representation of a floating point number may not be unique:
- For example, the number 13.25 can be represented as
$$1101.01_2 \cdot (2^0) = 110.101_2 \cdot (2^1) = 11.0101_2 \cdot (2^2) = 1.10101_2 \cdot (2^3)$$
- A floating point number is normalized when the integer part of its mantissa is forced to be exactly 1 and its fraction is adjusted accordingly.
- The leading 1 is **implicit**. It is not part of the 32 bits number.
- $1.F = 1. [2^{-1} \ 2^{-2} \dots 2^{-23}]$ .

# IEEE 754 standard

## Example 1

Represent  $3215.020002_{10}$

Decimal Value Entered: 3215.020002

Single precision (32 bits):

Binary: Status: normal

Bit 31 Sign Bit	Bits 30 - 23 Exponent Field	Bits 22 - 0 Significand
0	100 0101 0	1 .100 1000 1111 0000 0101 0010
0: + 1: -	Decimal value of exponent field and exponent 138 - 127 = 11	Decimal value of the significand 1.5698340

Hexadecimal: 4548F052      Decimal: 3215.0200

<http://babbage.cs.qc.cuny.edu/IEEE-754.old/Decimal.html>



Represent  $3215.020002_{10} \times 2 = 6430.040004_{10}$

Decimal Value Entered:

Single precision (32 bits):

Binary: Status:

Bit 31 Sign Bit	Bits 30 - 23 Exponent Field	Bits 22 - 0 Significand
<input type="text" value="0"/>	<input type="text" value="10001011"/>	<input type="text" value="1.10010001111000001010010"/>
0: + 1: -	Decimal value of exponent field and exponent <input type="text" value="139"/> - 127 = <input type="text" value="12"/>	Decimal value of the significand <input type="text" value="1.5698340"/>

Hexadecimal:  Decimal:

Represent  $3215.020002_{10}/4 = 803.7550005_{10}$

Decimal Value Entered:

Single precision (32 bits):

Binary: Status:

Bit 31 Sign Bit	Bits 30 - 23 Exponent Field	Bits 22 - 0 Significand
<input type="text" value="0"/>	<input type="text" value="10001000"/>	<input type="text" value="1.10010001111000001010010"/>
0: + 1: -	Decimal value of exponent field and exponent <input type="text" value="136"/> - 127 = <input type="text" value="9"/>	Decimal value of the significand <input type="text" value="1.5698340"/>

Hexadecimal:  Decimal:

- To multiply and divide by 2 is easy using floating-point numbers.
- Not every real number can be represented with floating-point format.
- Floating-point numbers are auto range !

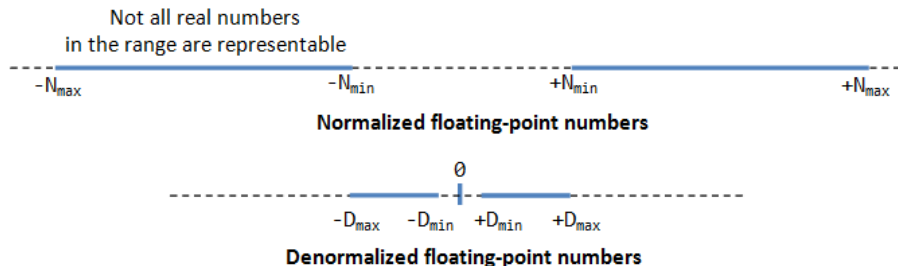
# IEEE 754 standard

Why auto range?

<b>2<sup>E</sup></b>	<b>MIN F</b> 1..000000000000000000000000 (b2) 1 (b10)	<b>MAX F</b> 1..111111111111111111111111 (b2) 1.999999881 (b10)
2 <sup>-126</sup>	1.1755E-38	2.3510E-38
2 <sup>-30</sup>	9.3132E-10	1.8626E-09
2 <sup>-20</sup>	9.5367E-07	1.9073E-06
2 <sup>-10</sup>	9.7656E-04	1.9531E-03
2 <sup>-3</sup>	0.12500000	0.24999999
2 <sup>-2</sup>	0.25000000	0.49999997
2 <sup>-1</sup>	0.50000000	0.99999994
<b>2<sup>0</sup></b>	<b>1.00000000</b>	<b>1.99999988</b>
2 <sup>1</sup>	2.00000000	3.99999976
2 <sup>2</sup>	4.00000000	7.99999952
2 <sup>3</sup>	8.00000000	15.99999905
2 <sup>10</sup>	1.0240E+03	2.0480E+03
2 <sup>20</sup>	1.0486E+06	2.0972E+06
2 <sup>30</sup>	1.0737E+09	2.1475E+09
2 <sup>127</sup>	1.7014E+38	3.4028E+38

# IEEE 754 standard

## De-normalized Form



- Normalized form has a serious problem.
- The number zero cannot be represent with an implicit leading 1!
- De-normalized form is devised to represent zero and small numbers.
- $E = 0 \Rightarrow E - bias = -127$
- Implicit leading  $0.F = 0.$   $[2^{-1} \ 2^{-2} \dots 2^{-23}]$ .

Represent  $-3.4\text{E-}39_{10}$ Decimal Value Entered: Single precision (32 bits):Binary: Status: 

Bit 31 Sign Bit	Bits 30 - 23 Exponent Field	Bits 22 - 0 Significand
<input type="text" value="1"/>	<input type="text" value="00000000"/>	<input type="text" value="0.1001010000010111010001"/>
0: + 1: -	Decimal value of exponent field and exponent <input type="text" value="0"/> - 127 = <input type="text" value="-127"/>	Decimal value of the significand <input type="text" value="0.5784800"/>

Hexadecimal:  Decimal:

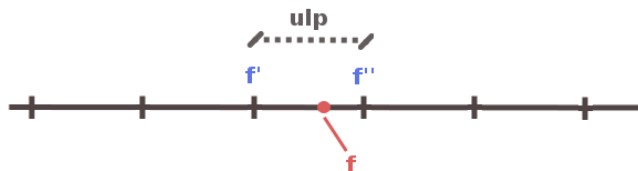
- **Zero:**  $E = 0, F = 0$ . Two representations: **+0** ( $S = 0$ ) and **-0** ( $S = 1$ ).
- **Inf** (Infinity):  $E = 0xFF, F = 0$ . Two representations: **+Inf** ( $S = 0$ ) and **-Inf** ( $S = 1$ ).
- **NaN** (Not a Number):  $E = 0xFF, F \neq 0$ . A value that cannot be represented as a real number (e.g.  $0/0$ ).

### MATLAB

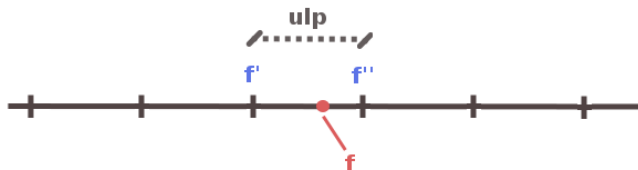
```
1 » a = 1/0
2 » a = Inf
3 » b = exp(1000)
4 » b = Inf
5 » c = log(0)
6 » c = -Inf
```

### MATLAB

```
1 » d = -1/0
2 » d = -Inf
3 » e = 0/0
4 » d = NaN
5 » f = Inf/Inf
6 » d = NaN
```



- $ulp$  (unit of least precision). In MATLAB, `eps()`.
- $f$ , significant,  $f = 1.F$ .
- $f'$  and  $f''$  being two successive multiples of  $ulp$ .
- Assume that  $f' < f < f''$ .
- $f'' = f' + ulp$ .
- Then, the rounding function  $round(f)$  associates to  $f$  either  $f'$  or  $f''$ , according to some rounding strategy.



Rounding schemes are:

- 1 **Truncation** (also called *round toward 0* or *chopping*):
  - if  $f$  is positive,  $\text{round}(f) = f'$ .
  - if  $f$  is negative,  $\text{round}(-f) = f''$ .
- 2 **Round toward plus infinity**:  $\text{round}(f) = f''$ .
- 3 **Round toward minus infinity**:  $\text{round}(f) = f'$ .
- 4 **Round to nearest** (default):
  - if  $f < f' + \text{ulp}/2$ ,  $\text{round}(f) = f'$ .
  - if  $f \geq f' + \text{ulp}/2$ ,  $\text{round}(f) = f''$ .



Dynamic range is defined as,

$$DR_{db} = 20 \log_{10} \left( \frac{\text{largest possible word value}}{\text{smallest possible word value}} \right) \quad [\text{dB}]$$

Dynamic range for floating-point numbers is defined as,

$$DR_{dB} \approx 6.02 \cdot 2^{b_E}$$

where  $b_E$  is the number of bits of  $E$ .

For single precision (32-bits):

$$DR_{dB} \approx 6.02 \cdot 2^8 \approx 1541 \text{ dB}$$

For 32-bits fixed point:

$$DR_{dB} \approx 6.02 \cdot 31 \approx 186 \text{ dB}$$

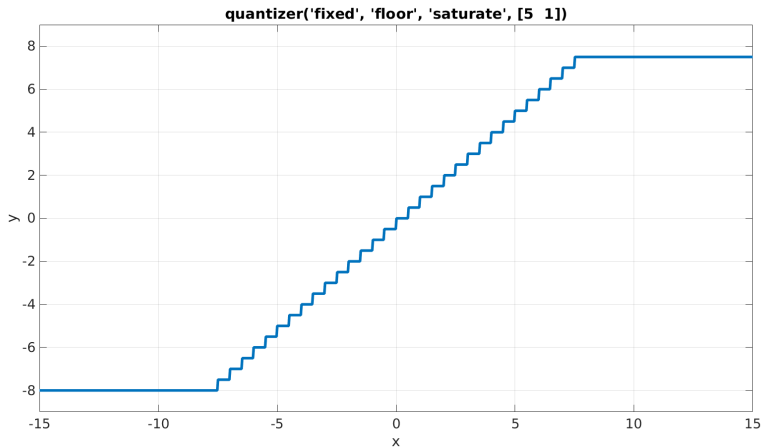
# Precision

## Fixed-point precision

- Precision is constant throughout all fixed-point numbers' range.
- Precision is  $2^{-n}$ .

## MATLAB

```
1 » % Fixed-point quantizer
2 » q = quantizer('fixed','floor','saturate',[5 1]);
3 % [wordlength fractionlength]
4 » u = linspace(-15,15,1000);
5 » y1 = quantize(q,u);
6 » plot(u,y1); title(tostring(q))
```



# Precision

## Floating-point precision

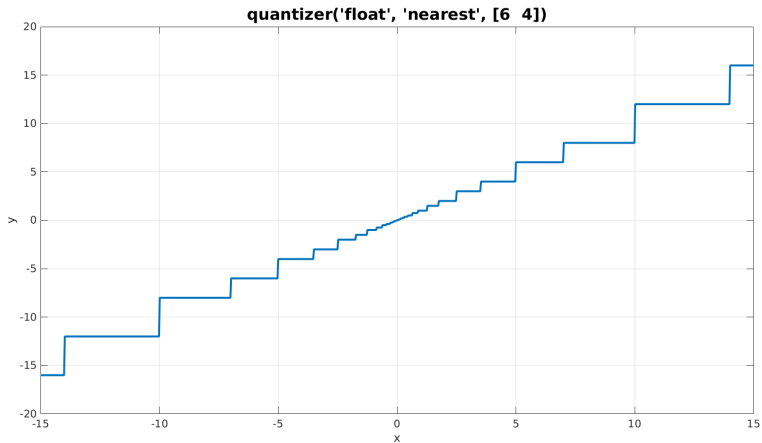
- Precision is **not** constant throughout all floating-point numbers' range.
- As the numbers get larger, the precision gets larger as well.
- Precision is  $2^E \cdot 2^{-23}$  for single precision.

## MATLAB

```
1 » % Floating-point quantizer
2 » q = quantizer([6 4], 'float', 'nearest');
3 % [wordlength exponentlength]
4 » y2 = quantize(q,u);
5 » plot(u,y2); title(tostring(q))
```

# Precision

## Floating-point precision



# Precision

Precision is not constant

<b>2<sup>E</sup></b>	<b>MIN F</b> 1..000000000000000000000000 (b2) 1 (b10)	<b>MAX F</b> 1..11111111111111111111111111 (b2) 1.999999881 (b10)	<b>PRECISION</b> 2 <sup>E</sup> * 2 <sup>-23</sup>
2 <sup>-126</sup>	1.1755E-38	2.3510E-38	1.4013E-45
2 <sup>-30</sup>	9.3132E-10	1.8626E-09	1.1102E-16
2 <sup>-20</sup>	9.5367E-07	1.9073E-06	1.1369E-13
2 <sup>-10</sup>	9.7656E-04	1.9531E-03	1.1642E-10
2 <sup>-3</sup>	0.12500000	0.24999999	1.4901E-08
2 <sup>-2</sup>	0.25000000	0.49999997	
2 <sup>-1</sup>	0.50000000	0.99999994	5.9605E-08
<b>2<sup>0</sup></b>	<b>1.00000000</b>	<b>1.99999988</b>	<b>1.1921E-07</b>
2 <sup>1</sup>	2.00000000	3.99999976	2.3842E-07
2 <sup>2</sup>	4.00000000	7.99999952	4.7684E-07
2 <sup>3</sup>	8.00000000	15.99999905	9.5367E-07
2 <sup>10</sup>	1.0240E+03	2.0480E+03	1.2207E-04
2 <sup>20</sup>	1.0486E+06	2.0972E+06	1.2500E-01
2 <sup>30</sup>	1.0737E+09	2.1475E+09	1.2800E+02
2 <sup>127</sup>	1.7014E+38	3.4028E+38	2.0282E+31

When calculations involve large and small numbers at the same time, the loss of precision affects the small number and the result.

### MATLAB

```
1 » (2^53 + 1) - 2^53
2 » ans = 0
3 » x = 0;
4 » t = tan(x) - sin(x)/cos(x)
5 » t = 0
6 » x = 1;
7 » t = tan(x) - sin(x)/cos(x)
8 » t = 2.2204e-16 % eps(1)
9 » if (t == 0)
10 »     disp('Start nuclear fusion')
11 » else
12 »     disp('Do not start nuclear fusion')
```

Perform  $0.5 + (-0.4375)$  using 4 bits for the mantissa.

$$0.5_{10} = 0.1000_2 \times 2^0 = 1.0000_2 \times 2^{-1} \text{ (normalised)}$$

$$-0.4375_{10} = -0.0111_2 \times 2^0 = -1.1100_2 \times 2^{-2} \text{ (normalised)}$$

- 1 Match exponents to the bigger one.  
Apply  $n$  right shifts to  $-0.4375$  where  $n = (\text{exponent1} - \text{exponent2}) = (-1 + 2) = 1$ .  
 $-0.4375 = -1.1100_2 \times 2^{-2} = -\mathbf{0.1110_2} \times \mathbf{2^{-1}}$
- 2 Add the mantissas.  
 $(1.0000_2 - 0.1110_2) \times 2^{-1} = 0.0010_2 \times 2^{-1}$
- 3 Normalise the sum, checking for overflow/underflow:  
 $0.0010_2 \times 2^{-1} = 1.0000_2 \times 2^{-4} = \mathbf{0.0625}$   
 $-126 \leq -4 \leq 127$ , no overflow or underflow
- 4 Round the sum.  
The sum fits in 4 bits so rounding is not required



# Precision

## Sum of not similar floating-point numbers

Perform  $1e10 + 1500$  using IEEE-754 single precision.

$$10,000,000,000 = -1.00101010000001011111001_2 \times 2^{33} \text{ (normalised)}$$

$$1500 = -1.01110111000000000000000_2 \times 2^{10} \text{ (normalised)}$$

- 1 Match exponents to the bigger one.

Apply  $n$  right shifts to 1500 where  $n = (\text{exponent1} - \text{exponent2}) = (33 - 10) = 23$

$$1500 \approx 0.000000000000000000000001_2 \times 2^{33} = (1.0 \times 2^{-23}) \times 2^{33} = 1.0 \times 2^{10} = 1024$$

- 2 Add the mantissas.

$$(1.00101010000001011111001_2 + 0.000000000000000000000001_2) \times 2^{33}$$

- 3 Normalise the sum, checking for overflow/underflow:

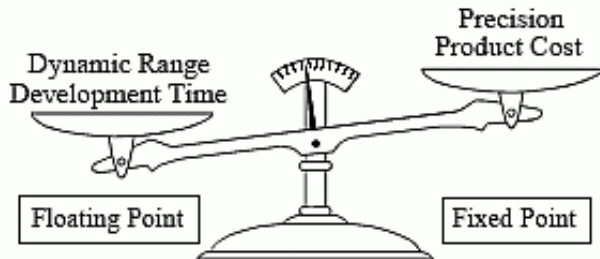
$$1.010101000000101111010_2 \times 2^{33} = 1.16415333710 \times 2^{33} = 10,000,001,024$$

$-126 \leq 33 \leq 127$ , no overflow or underflow

- 4 Round the sum.

The sum fits in 23 bits so rounding is not required

## Fixed-point vs floating-point



- 1 Jean-Pierre Deschamps, Gustavo D. Sutter, and Enrique Cantó. Guide to FPGA Implementation of Arithmetic Functions, Chapter 12.